

## Bellabeat Case Study

I have already used SQL and pivot tables in spreadsheets to do some initial analysis. I want to rely on R programming language to help create visualizations of the data. I will begin by loading the appropriate packages.

```
install.packages("tidyverse")

## Installing package into '/cloud/lib/x86_64-pc-linux-gnu-library/4.2'
## (as 'lib' is unspecified)

library(tidyverse)

## -- Attaching packages ----- tidyverse 1.3.1 --
## v ggplot2 3.3.5      v purrr   0.3.4
## v tibble  3.1.6      v dplyr  1.0.9
## v tidyr   1.2.0      v stringr 1.4.0
## v readr   2.1.2      v forcats 0.5.1

## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()

library(readr)
```

I will now upload my datasets. Because I was able to combine the hourly data in spreadsheets, I can upload that on its own, but I am going to have to rely on joins to aggregate the sleep data into the daily dataset.

```
dailyactivity <- read_csv("Bellabeat_Case_Study/Bellabeat_daily.csv")

## Rows: 940 Columns: 7
## -- Column specification -----
## Delimiter: ","
## chr (1): ActivityDate
## dbl (6): Id, TotalSteps, TotalDistance, TotalActiveMinutes, SedentaryMinutes...
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.

dailysleep <- read_csv("Bellabeat_Case_Study/Bellabeat_sleep.csv")

## Rows: 413 Columns: 4
## -- Column specification -----
## Delimiter: ","
## chr (1): SleepDate
## dbl (3): Id, TotalMinutesAsleep, TotalTimeInBed
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.

hourly <- read_csv("Bellabeat_Case_Study/Bellabeat_hourly.csv")

## Rows: 22099 Columns: 6
```

```
## -- Column specification -----
## Delimiter: ","
## chr (1): Activity Day
## dbl (4): Id, Calories, Total Intensities, Steps
## time (1): Activity Hour
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
weekdaysleep <- read_csv("Bellabeat_Case_Study/Bellabeat_avgSleepDay.csv")
```

```
## Rows: 504 Columns: 7
## -- Column specification -----
## Delimiter: ","
## chr (3): Weekday, ActivityDate, Avg_Distance
## dbl (4): Avg_Sedentary, Avg_Sleep, Avg_Steps, Avg_Calories
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
weekday <- read_csv("Bellabeat_Case_Study/Bellabeat_weekday.csv")
```

```
## Rows: 7 Columns: 5
## -- Column specification -----
## Delimiter: ","
## chr (1): Weekday
## dbl (4): Sedentary, Sleep, Steps, Calories
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
head(dailyactivity)
```

```
## # A tibble: 6 x 7
##       Id ActivityDate TotalSteps TotalDistance TotalActiveMinu~ SedentaryMinutes
##   <dbl> <chr>          <dbl>         <dbl>          <dbl>          <dbl>
## 1 1.50e9 4/12/2016        13162          8.5            366            728
## 2 1.50e9 4/13/2016        10735          6.97           257            776
## 3 1.50e9 4/14/2016        10460          6.74           222           1218
## 4 1.50e9 4/15/2016         9762          6.28           272            726
## 5 1.50e9 4/16/2016       12669          8.16           267            773
## 6 1.50e9 4/17/2016         9705          6.48           222            539
## # ... with 1 more variable: Calories <dbl>
```

```
head(dailysleep)
```

```
## # A tibble: 6 x 4
##       Id SleepDate TotalMinutesAsleep TotalTimeInBed
##   <dbl> <chr>          <dbl>         <dbl>
## 1 1503960366 4/12/2016         327            346
## 2 1503960366 4/13/2016         384            407
## 3 1503960366 4/15/2016         412            442
## 4 1503960366 4/16/2016         340            367
## 5 1503960366 4/17/2016         700            712
## 6 1503960366 4/19/2016         304            320
```

```
head(hourly)
```

```
## # A tibble: 6 x 6
##       Id `Activity Day` `Activity Hour` Calories `Total Intensities` Steps
##       <dbl> <chr>         <time>         <dbl>         <dbl> <dbl>
## 1 1503960366 4/12/2016      00:00             81             20    373
## 2 1503960366 4/12/2016      01:00             61              8    160
## 3 1503960366 4/12/2016      02:00             59              7    151
## 4 1503960366 4/12/2016      03:00             47              0     0
## 5 1503960366 4/12/2016      04:00             48              0     0
## 6 1503960366 4/12/2016      05:00             48              0     0
```

```
head(weekdaysleep)
```

```
## # A tibble: 6 x 7
##   Weekday ActivityDate Avg_Sedentary Avg_Sleep Avg_Steps Avg_Calories
##   <chr>    <chr>         <dbl>    <dbl>    <dbl>    <dbl>
## 1 Thursday 5/5/2016          749.    362.    10255.    2550.
## 2 Tuesday 4/26/2016          780.    369.    9290.    2444.
## 3 Thursday 4/21/2016          791.    376.    9698.    2579.
## 4 Friday 4/29/2016          716.    386.    7910.    2269.
## 5 Saturday 4/16/2016          710.    392.    8615.    2494.
## 6 Friday 4/22/2016          739.    393.    8377.    2456.
## # ... with 1 more variable: Avg_Distance <chr>
```

```
head(weekday)
```

```
## # A tibble: 6 x 5
##   Weekday Sedentary Sleep Steps Calories
##   <chr>    <dbl> <dbl> <dbl>    <dbl>
## 1 Friday      741   405  7936    2332
## 2 Monday      718   419  9262    2435
## 3 Saturday    682   418  9857    2504
## 4 Sunday      686   455  7295    2279
## 5 Thursday    660   405  7810    2208
## 6 Tuesday     740   405  9183    2499
```

And now to join the daily data.

```
daily_join <- right_join(dailyactivity, dailysleep, by = c('Id'='Id', 'ActivityDate'='SleepDate'))
head(daily_join)
```

```
## # A tibble: 6 x 9
##       Id ActivityDate TotalSteps TotalDistance TotalActiveMinu~ SedentaryMinutes
##       <dbl> <chr>         <dbl>         <dbl>         <dbl>         <dbl>
## 1 1.50e9 4/12/2016      13162           8.5           366           728
## 2 1.50e9 4/13/2016      10735           6.97          257           776
## 3 1.50e9 4/15/2016       9762           6.28          272           726
## 4 1.50e9 4/16/2016      12669           8.16          267           773
## 5 1.50e9 4/17/2016       9705           6.48          222           539
## 6 1.50e9 4/19/2016      15506           9.88          345           775
## # ... with 3 more variables: Calories <dbl>, TotalMinutesAsleep <dbl>,
## #   TotalTimeInBed <dbl>
```


Now we have a complete, clean dataset for the daily data. We are now going to try out some different visualizations.

The first question became, “What do we qualify as a healthy user?” After some research, these are the following metrics we are going to use *Even though gender is not specified, I’m going to use suggested metrics related to women* - Recommended for adults to get between 7-9 hours of sleep a night - Recommended for adults to get around 10,000 steps a day - Recommended for adult women to burn around 2,000 calories a day

I will use these to qualify a “healthy” user.

Using spreadsheets and SQL I have already discovered: - 33 users provided data - 24 users provided sleep data - 413 individual sleeps recorded - Only 119 sleep days that fall between 7-9 hours - 19 out of the 24 users averaged suggested sleep over the 30 days

```
ggplot(weekday, mapping = aes(x=reorder(Weekday,+Sleep), y=Sleep)) + geom_col()
```




This first visualization shows the average sleep grouped by weekday. So on average, Fridays, Thursdays, and Tuesdays get the least amount of sleep.

So the next question is how do we help users get more sleep? My hypothesis is that increased activity leads to increased sleep.


```
ggplot(daily_join, mapping = aes(x=TotalMinutesAsleep,y=TotalActiveMinutes)) + geom_smooth()
```

```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```

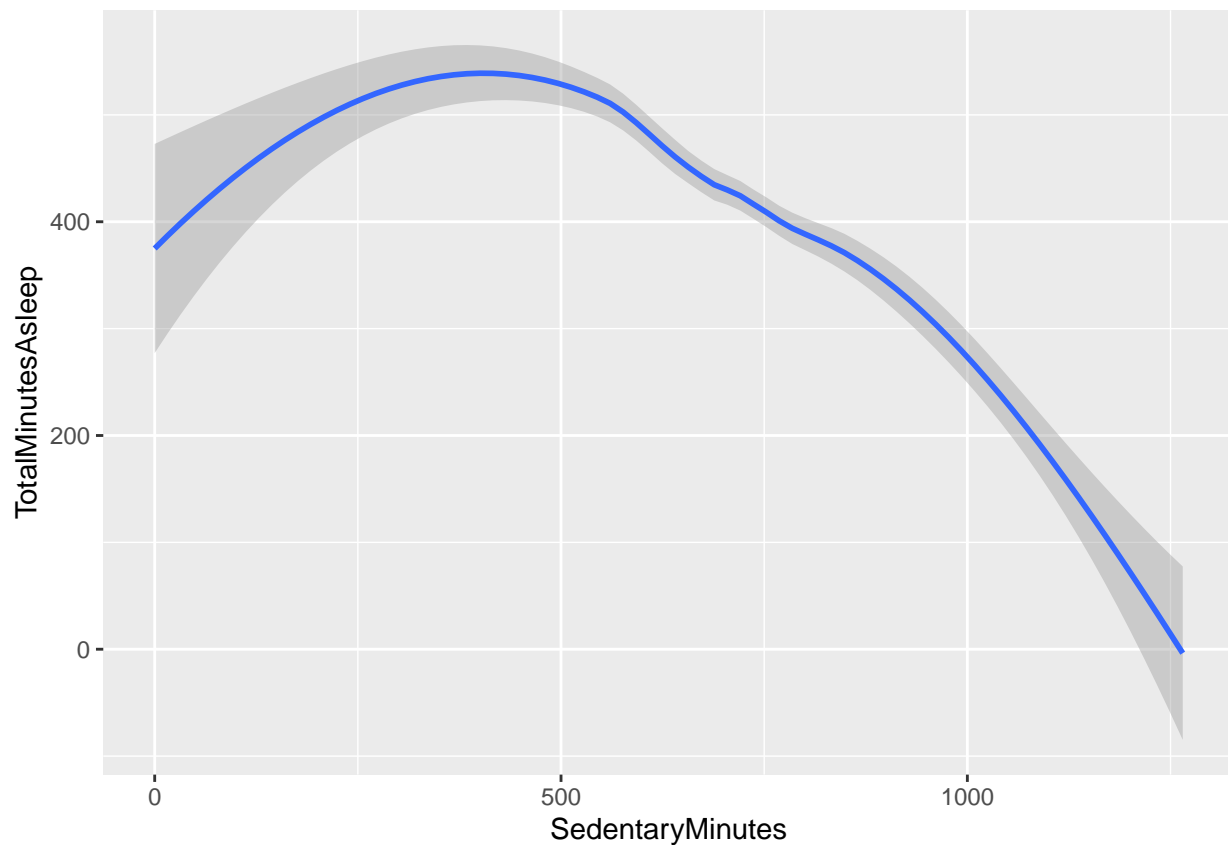


There actually isn't a real correlation between the two. So using my remaining data points, I'm going to check for correlation between them

```
ggplot(daily_join, mapping = aes(x=TotalDistance,y=TotalMinutesAsleep)) + geom_smooth()
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```



```
ggplot(daily_join, mapping = aes(x=SedentaryMinutes,y=TotalMinutesAsleep)) + geom_smooth()  
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```



Of the three, the strongest negative correlation is between sedentary minutes and total minutes asleep. Let me try to look at this from another angle.

```
AvgId <- aggregate(cbind(Active <- daily_join$TotalActiveMinutes, Sedentary <- daily_join$SedentaryMinutes),
  AvgId <- setNames(AvgId, c("Id", "ActivityMin", "SedentaryMin", "Steps", "SleepMin"))
  arrange(AvgId, desc(SleepMin), .by_group = FALSE)
```


##	Id	ActivityMin	SedentaryMin	Steps	SleepMin
## 1	1844505072	147.3333	443.3333	3477.000	652.0000
## 2	2026352035	256.8929	653.9643	5618.679	506.1786
## 3	6117666160	364.5556	531.9444	8823.833	478.7778
## 4	4319703577	259.2308	642.6923	7125.423	476.6538
## 5	5553957443	242.6129	668.3548	8612.581	463.4839
## 6	7086361926	234.0000	723.6667	10290.500	453.1250
## 7	6962181067	287.1290	662.3226	9794.806	448.0000
## 8	2347167796	271.2000	628.4000	8533.200	446.8000
## 9	8378563200	226.5625	715.3750	8832.938	443.3438
## 10	8792009665	178.4000	807.8000	3443.267	435.6667
## 11	5577150313	296.2308	667.3077	9260.077	432.0000
## 12	4702921684	292.1429	693.0357	9226.357	421.1429
## 13	1927972279	85.0000	977.2000	1490.000	417.0000
## 14	4388161847	286.7500	751.4583	10974.708	403.1250
## 15	4445114986	217.9643	787.3214	4756.179	385.1786
## 16	1503960366	291.3200	759.2800	12405.680	360.2800
## 17	6775888955	107.0000	964.0000	3499.000	349.6667
## 18	4020332650	249.0000	841.8750	6596.750	349.3750
## 19	8053475328	301.0000	837.3333	19078.667	297.0000
## 20	1644430081	263.2500	920.5000	7967.750	294.0000

```
## 21 3977333714    262.6429    716.2143 11218.000 293.6429
## 22 4558609924    313.0000    1028.4000 8139.000 127.6000
## 23 7007744171    220.0000    1148.5000 5115.500  68.5000
## 24 2320127002    242.0000    1129.0000 5079.000  61.0000
```

That gave me a tibble of the basic averages per each user with data on activity minutes, sedentary minutes, steps, and sleep minutes. I now want to see if the averages have any more correlation.

```
ggplot(AvgId, mapping = aes(x=SedentaryMin, y=SleepMin)) + geom_smooth()
```


```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```



```
ggplot(AvgId, mapping = aes(x=Steps, y=SleepMin)) + geom_smooth()
```


```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```





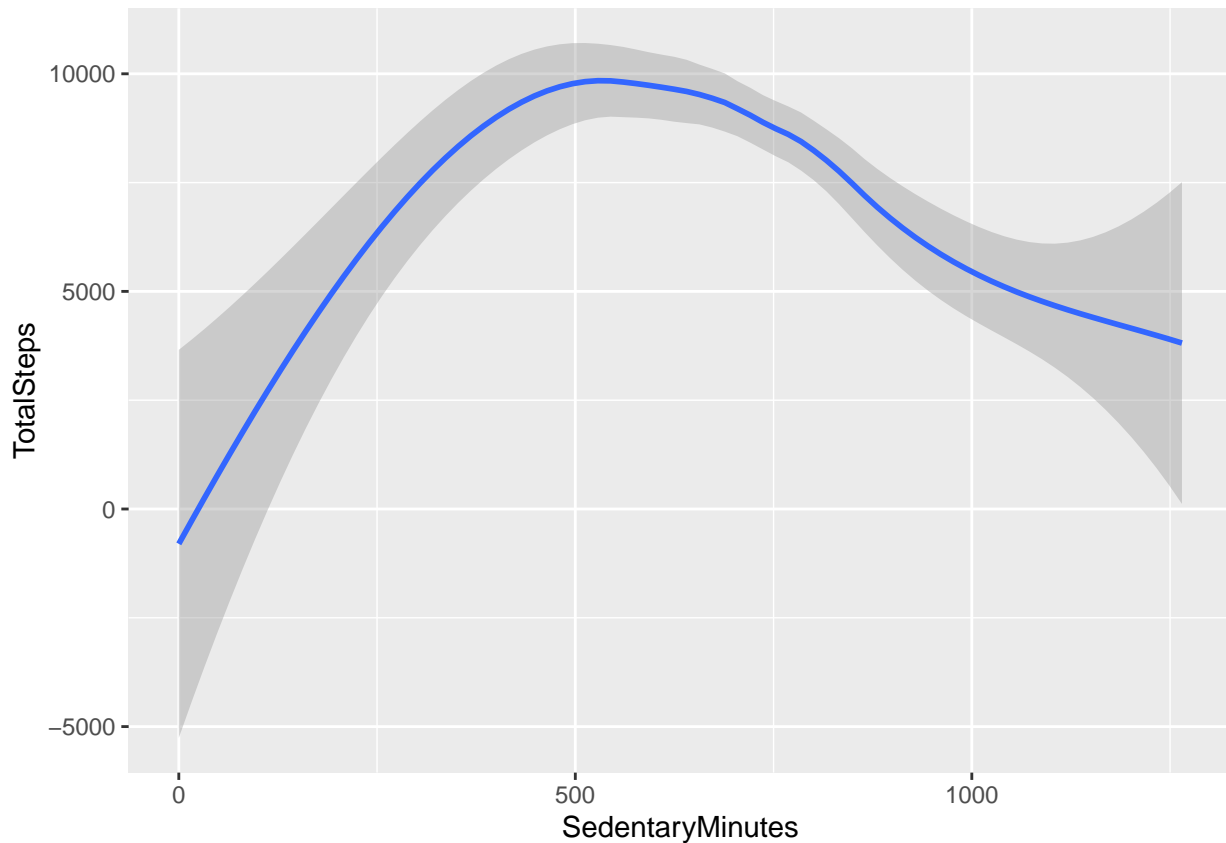
```
ggplot(AvgId, mapping = aes(x=ActivityMin, y=Steps)) + geom_smooth()
```

```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```



```
ggplot(daily_join, mapping = aes(x=SedentaryMinutes,y=TotalSteps)) +geom_smooth()
```

```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```



Even using a different set of averages, it still shows a negative correlation between time spent sedentary and time spent asleep. Rather than increasing exercise, Bellabeat should focus on reducing sedentary time to increase sleep.

```
hourly_mod <- aggregate(cbind(hourly$Calories, as.integer(hourly$`Total Intensities`), hourly$Steps), 1,
  setNames(hourly_mod, c("ActivityHour", "Calories", "Intensity", "Steps"))
```

##	ActivityHour	Calories	Intensity	Steps
## 1	00:00:00	71.80514	7.3158458	140.389722
## 2	01:00:00	70.16506	3.5069668	50.526259
## 3	02:00:00	69.18650	1.8510182	31.326902
## 4	03:00:00	67.53805	0.9721329	11.675241
## 5	04:00:00	68.26180	0.3798283	3.589056
## 6	05:00:00	81.70815	0.1673820	2.115880
## 7	06:00:00	86.99678	0.8979592	13.676692
## 8	07:00:00	94.47798	2.5273899	38.126745
## 9	08:00:00	103.33727	8.8818475	185.514501
## 10	09:00:00	106.14286	22.4167562	1063.537057
## 11	10:00:00	110.46071	15.9149623	443.417653
## 12	11:00:00	109.80690	14.1974110	462.752967
## 13	12:00:00	117.19740	18.4446855	559.980477
## 14	13:00:00	115.30945	21.2931596	708.247557
## 15	14:00:00	115.73290	23.2877307	737.046688
## 16	15:00:00	106.63716	14.5890710	450.953005
## 17	16:00:00	113.32745	18.9834620	568.424476
## 18	17:00:00	122.75276	20.8410596	698.151214
## 19	18:00:00	123.49227	43.6545254	1520.247241
## 20	19:00:00	121.48455	31.5673289	984.312362

```
## 21      20:00:00 102.35762 31.6004415 1085.163355
## 22      21:00:00  96.05635 37.0099448 1301.507182
## 23      22:00:00  88.26549 20.8440265  656.394912
## 24      23:00:00  77.59358 24.0365449  745.544850
```

```
arrange(hourly_mod, desc("Intensity"), .by_group = TRUE)
```

```
##      Group.1      V1      V2      V3
## 1 00:00:00 71.80514 7.3158458 140.389722
## 2 01:00:00 70.16506 3.5069668  50.526259
## 3 02:00:00 69.18650 1.8510182  31.326902
## 4 03:00:00 67.53805 0.9721329  11.675241
## 5 04:00:00 68.26180 0.3798283   3.589056
## 6 05:00:00 81.70815 0.1673820   2.115880
## 7 06:00:00 86.99678 0.8979592  13.676692
## 8 07:00:00 94.47798 2.5273899  38.126745
## 9 08:00:00 103.33727 8.8818475 185.514501
## 10 09:00:00 106.14286 22.4167562 1063.537057
## 11 10:00:00 110.46071 15.9149623 443.417653
## 12 11:00:00 109.80690 14.1974110 462.752967
## 13 12:00:00 117.19740 18.4446855 559.980477
## 14 13:00:00 115.30945 21.2931596 708.247557
## 15 14:00:00 115.73290 23.2877307 737.046688
## 16 15:00:00 106.63716 14.5890710 450.953005
## 17 16:00:00 113.32745 18.9834620 568.424476
## 18 17:00:00 122.75276 20.8410596 698.151214
## 19 18:00:00 123.49227 43.6545254 1520.247241
## 20 19:00:00 121.48455 31.5673289  984.312362
## 21 20:00:00 102.35762 31.6004415 1085.163355
## 22 21:00:00  96.05635 37.0099448 1301.507182
## 23 22:00:00  88.26549 20.8440265  656.394912
## 24 23:00:00  77.59358 24.0365449  745.544850
```

So this now breaks down average calories, intensities, and steps based on hour of the day. Although I had issues arranging the data set, you can still see that (excluding normal sleeping hours) 5AM-8AM have the lowest intensities recorded


```
weekdaysleep %>%
  group_by(Weekday)%>%
  summarise(sedentary = mean(Avg_Sedentary, na.rm=FALSE), sleep = mean(Avg_Sleep, na.rm=FALSE))
```

```
## # A tibble: 8 x 3
##   Weekday  sedentary sleep
##   <chr>      <dbl> <dbl>
## 1 Friday      741.  405.
## 2 Monday      718.  419.
## 3 Saturday    682.  418.
## 4 Sunday      686.  455.
## 5 Thursday    660.  405.
## 6 Tuesday     740.  405.
## 7 Wednesday   712.  435.
## 8 <NA>         NA    NA
```

This breaks down daily averages of sedentary time and sleep time.

```
ggplot(weekdaysleep, mapping = aes(x=Weekday, y=Avg_Sedentary)) + geom_col()
```

## Warning: Removed 473 rows containing missing values (position\_stack).



As shown here, Tuesday-Thursday have the highest average sedentary time.